


## ✓ nmi | spring 2024

### special lecture 04 : neural networks, briefly

#### ✓ 1 perceptron

##### ✓ 1.1 review



##### ✓ code, perceptron

```
1 import numpy as np
2
3 class Perceptron:
4
5     def __init__(self, learning_rate=0.01, epochs=1000, weights=None, bias=None, \
6                   tol=1e-8, activation_func="step", delta_func=None, leaky=0.1):
7         self.learning_rate = learning_rate
8         self.epochs = epochs
9         self.weights = weights
10        self.bias = bias
11        self.tol = tol
12        self.delta_func = delta_func
13        self.lr_leaky = leaky
14
15        if activation_func is None: ...
17        if isinstance(activation_func, str): ...
33        else: ...
35
36    def copy(self): ...
49
50    def fit(self, X, y, display=False):
51        n_samples, n_features = X.shape
52
53        # init parms
54        if self.weights is None or self.bias is None: ...
62
63        # learn
64        rmse_old = self.tol*2
65        for i in range(self.epochs):
66            errors = 0
67            for j in range(n_samples):
68                # forward pass
69                z = np.dot(X[j], self.weights) + self.bias
70                y_hat = self.activation_func(z)
71                error = y[j] - y_hat
72                errors += error**2
73                # backward pass
74                delta = self.delta_func(error, y_hat)
75                self.weights += self.learning_rate * X[j] * delta
76                self.bias += self.learning_rate * delta
77            rmse = np.sqrt(errors / n_samples)
78            if display:
79                print(f"Epoch {i}: rmse = {rmse}")
80            if rmse < self.tol:
81                break
82        return self
```

```

67     for x,target in zip(X,y):
68         z = np.dot(x,self.weights) + self.bias # linear_output
69         y_predicted = self.activation_func(z)
70
71         # update weights, bias
72         error = y_predicted - target
73         self.bias -= self.learning_rate*error
74         delta = self.delta_func(x,z)
75         self.weights -= self.learning_rate*error*delta*x
76
77         errors += pow(error,2)
78         rmse = np.sqrt(errors/len(y))
79         if (rmse < self.tol) or (rmse == rmse_old) or (i % 100 == 0):...
85         rmse_old = rmse
86
87     def predict(self,X):
88         linear_output = np.dot(X,self.weights) + self.bias
89         y_predicted = self.activation_func(linear_output)
90         return y_predicted
91
92     # functions: activation, delta
93     def leaky(self,x,alpha=None):...
97     def dleaky(self,x,z,alpha=None):...
101    def ReLU(self,x):...
103    def dReLU(self,x,z):...
105    def sigmoid(self,x):...
107    def dsigmoid(self,x,z):...
109    def step(self,x):...
111    def dstep(self,x,z):...
113    def tanh(self,x):...
115    def dtanh(self,x,z):...
117
118    # misc
119    def sign(self,x):...
121
1
2  if __name__ == "__main__":
3
4      # imports
5      import matplotlib.pyplot as plt
6      import numpy as np
7      import scipy as sp
8      from sklearn.model_selection import train_test_split
9      from sklearn import datasets
10     import statistics as st
11     from tabulate import tabulate
12
13     def accuracy(y_pred,y_true,tol=0):
14         if tol == 0:...
15         else:
16             accuracy = np.sum(abs(y_pred-y_true)<=tol) / len(y_true)
17         return accuracy
18
19     # init runtime
20     debug = False
21
22     # datasets from sklearn
23     X,y = datasets.make_blobs(n_samples=150,n_features=2,centers=2,cluster_std=1.05,random_state=2)
24     X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=123)
25     x0 = np.array([np.amin(X_train[:,0]),np.amax(X_train[:,0])]) # interval, training data
26
27     # plot, init
28     fig = plt.figure()
29     ax = fig.add_subplot(1,1,1)
30
31     if True: # plot data, training...
32
33     # weights, bias for everyone
34     n_features = X.shape[1]
35     if False: # suggested weights when random...
36     else: # suggested for ReLU...
37     weights = np.random.uniform(urng_from,urng_thru,n_features)
38     bias = np.random.uniform(urng_from,urng_thru)

```

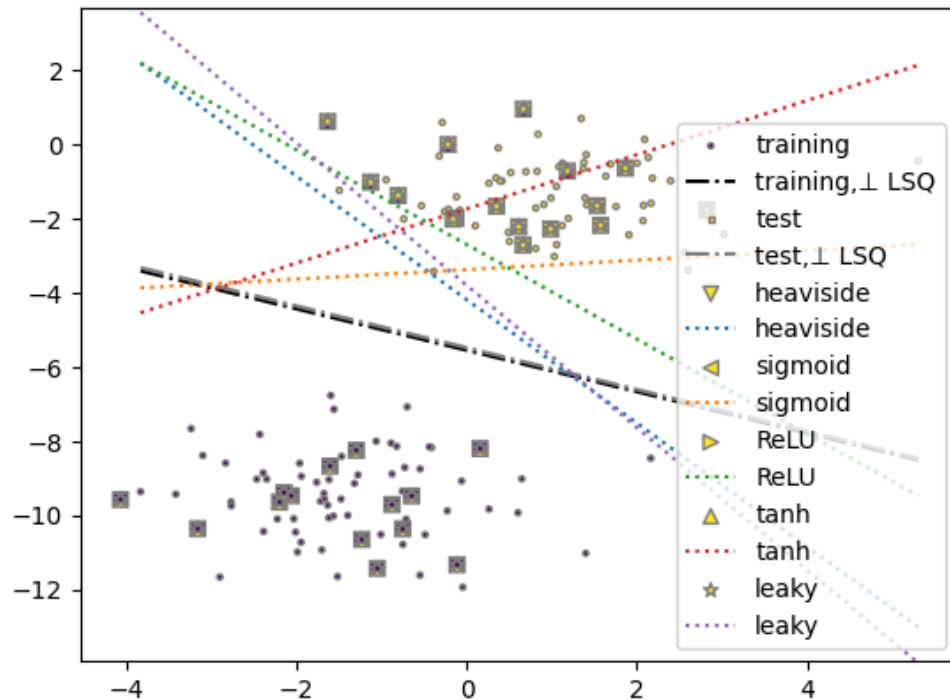
```

54
55 # init nn
56 ss_afname = ["heaviside","sigmoid","ReLU","tanh","leaky"]
57 ss_marker = ["v","<",">","^","*"] # markers for plot
58 data = [] # array for tabulate
59
60 # train and test per activation function
61 for s_afname,s_marker in zip(ss_afname,ss_marker):
62     p = Perceptron(learning_rate=0.005,epochs=1000,weights=weights.copy(),bias=bias,activation_func=s_afname)
63     if debug: # spacer...
64         p.fit(X_train,y_train,display=debug)
65
66     # prediction, classification
67     predictions = p.predict(X_test)
68     if debug: # accuracy wrt test data...
69         if debug: # accuracy wrt training data...
70
71     # plot, predictions
72     plt.scatter(X_test[:,0],X_test[:,1],label=s_afname,marker=s_marker,c=predictions,edgecolor="grey")
73
74     # plot, linear separation
75     x1 = -(p.weights[0]*x0+p.bias)/p.weights[1]
76     ax.plot(x0,x1,label=s_afname,ls=":") # hyperplane ~ decision boundary
77
78     # tabulate weights, bias
79     s_weights = ",".join(f"{weight:.8f}" for weight in p.weights)
80     s_accuracy = f"{accuracy(predictions,y_test,p.tol)*100:3.0f}"
81     if debug: ...
82     else: ...
83
84 # tabulate, show
85 print()
86 if debug: ...
87 else: ...
88
89 # plot, show
90 ymin = np.amin(X_train[:,1])
91 ymax = np.amax(X_train[:,1])
92 ax.set_ylim([ymin-2,ymax+3])
93 plt.title("\ndont look too close\n")
94 plt.legend(loc="lower right")
95 plt.show()
96

```

| af        | w                       | b           | test(%) |
|-----------|-------------------------|-------------|---------|
| heaviside | 0.04536806, 0.02725385  | 0.11406480  | 100     |
| sigmoid   | -0.01231378, 0.09611858 | 0.32406480  | 100     |
| ReLU      | 0.12265729, 0.09611858  | 0.25906480  | 100     |
| tanh      | 0.19444766, -0.26616686 | -0.46093520 | 33      |
| leaky     | 0.09304901, 0.04847492  | 0.18406480  | 100     |

dont look too close



## ✓ 1.2 predict()

forward propagation:

1. application of weights;
2. summation;
3. activation.

## ✓ 1.2.1 activation functions, revisited

## ✓ common activation functions

- linear. yes, thats right; it doesnt add a thing. think of it as a placeholder.

- heaviside: binary step. for beginners = my first activation function.

$$f(x) = \begin{cases} 0 & x < 0 \\ 1 & x \geq 0 \end{cases}$$

- sigmoid: input scaled  $x \mapsto [0, 1]$ . used for binary classification where output can be interpreted as a probability. a popular flavor.

$$f(x) = \frac{1}{1 + e^{-x}}$$

- hyperbolic tangent, tanh: input scaled  $x \mapsto [-1, 1]$  centered around zero. used in hidden layers to balance input signal.

$$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$$

- rectified linear unit (ReLU): blocks negative input. simple trick as an option when training deep neural networks. most popular for hidden layers, so they say.

$$f(x) = \max(0, x)$$

- leaky ReLU: allows very small, non-zero gradient when negative input. keeps neurons alive when not actively firing.

$$f(x) = \max(\alpha x, x), \quad \alpha \ll 1$$

- exponential linear unit (ELU): smooths negative inputs. may help maintain mean activation close to zero; an alternative for ReLU.

$$f(x) = \begin{cases} x & x > 0 \\ \alpha(e^x - 1) & x \leq 0 \end{cases}$$


- normalized exponential (softmax): turns logits into probabilities. popular for multi-class classification, picking the most likely one.

$$P(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

## ✓ visual, comparative

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3
4 if __name__ == "__main__":
5
```

## some activation functions



### ✓ choice per location

- output layer: sigmoid, softmax
- hidden layers: ReLU, leaky ReLU, tanh.

### ✓ 1.2.2 wrt bias, threshold

consider the heaviside activation function,

$$\hat{y} = \begin{cases} 1 & \text{if } \sum_1^n x_i w_i \geq \theta \\ 0 & \text{if } \sum_1^n x_i w_i < \theta \end{cases} = \begin{cases} 1 & \text{if } \sum_1^n x_i w_i - \theta \geq 0 \\ 0 & \text{if } \sum_1^n x_i w_i - \theta < 0 \end{cases} = \begin{cases} 1 & \text{if } \sum_0^n x_i w_i \geq \theta \\ 0 & \text{if } \sum_0^n x_i w_i < \theta \end{cases} \quad \text{where } x_0 = 1, w_0 = -\theta$$

that makes threshold vs bias look somewhat ambiguous; however, threshold is the decision boundary for the activation function and bias is an offset to the weighted sum of inputs.

### ✓ 1.3 fit()

backpropagation ~ "backward propagation of errors":

1. forward pass to generate output;
2. backward pass where error informs changes in weights and bias.

the gradient of the loss function expressed with the chain rule of calculus reveals how much error each weight contributes to the error, providing a path for its adjustment.

gradient descent of the loss function wrt weights:

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \phi} \cdot \frac{\partial \phi}{\partial z} \cdot \frac{\partial z}{\partial w}$$

where  $L$  is the loss function,  $\phi$  is the activation function,  $z$  is the weighted, biased input and  $w, b$  are the weights and bias.

weights are moved in the direction of steepest decrease in loss and with magnitude adjusted by learning rate hyperparameter  $\eta$ . ie,

$$w_{new} = w_{old} - \eta \cdot \frac{\partial L}{\partial w}.$$

similarly for bias,

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \phi} \cdot \frac{\partial \phi}{\partial z} \cdot \frac{\partial z}{\partial b}$$

$$b_{new} = b_{old} - \eta \cdot \frac{\partial L}{\partial b}.$$

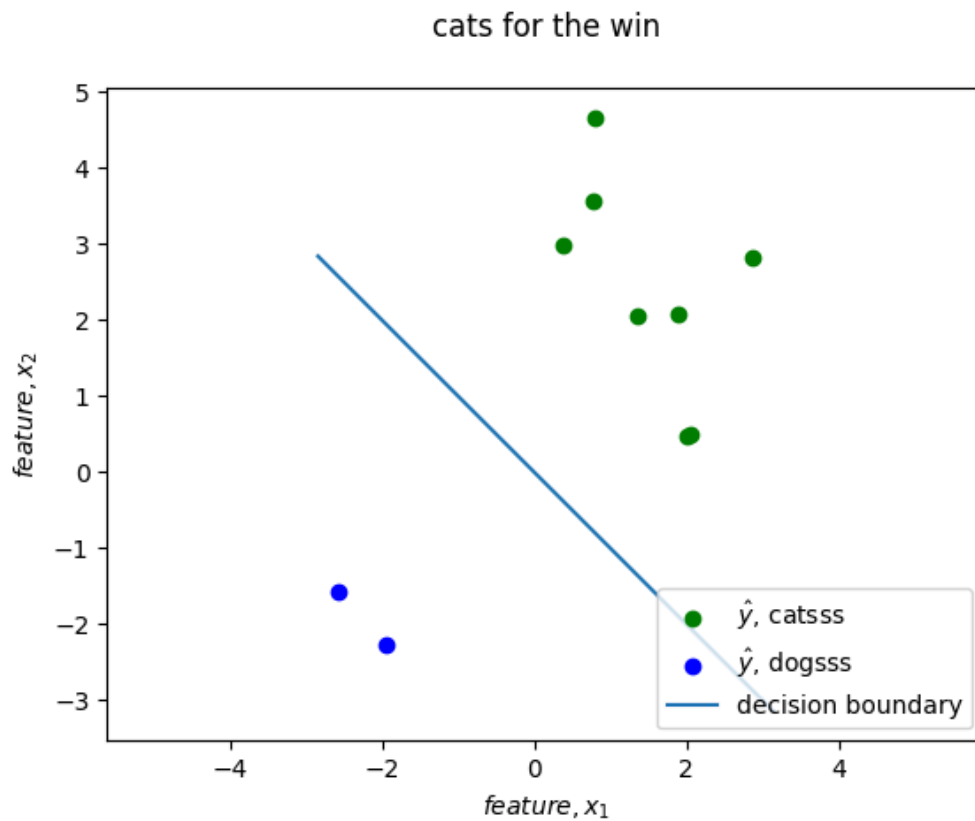
this repeats until model converges to a state where the loss is minimized and/or some other criteria are met.

### ✓ 1.3 linear separability

to reiterate, data sets are linearly separable if you can put a hyperplane between them.

### ✓ hyperplane, visual

### ✓ code, python



✓ perceptron convergence theorem

for any set of **linearly separable** labeled examples, the perceptron learning algorithm will halt after a **finite number of iterations**. ie, after a finite number of iterations, the algorithm **finds (not optimizes)** a vector that classifies perfectly all the examples. (mostly [mit](#))

✓ und so weiter

- **pure math** that allowed near neural nets: linear algebra, basic abstract algebra, set theory, functional analysis. so they say.

✓ theorems

- perceptron convergence theorem: rosenblatt [@princeton](#); novikoff [@oregonstate](#)
- perceptron capacity: cover counting theorem [@mit](#)

✓ bookshelf

- mit press, deep learning (textbook) [online](#)

... some elements of formal logic. [Github](#)